

# Fast Poisson Disk Sampling in Arbitrary Dimensions

Steven Yu-Chih Lin

National Taiwan University  
R04922170

Andrew Yu-Chian Wu

National Taiwan University  
R05922103

## ABSTRACT

In this project, we implemented a fast Poisson-Disk generating algorithm, introduced by Robert Bridson in 2007. This algorithm is a slightly revised version of dart throwing, and it generated samples in linear time, i.e.,  $O(N)$  complexity. The capability of being easily extended to and implemented in arbitrary dimension is one of the most important characteristics of this algorithm. We generated 4 sets of samples and demonstrated the results in our evaluation.

## INTRODUCTION

Sampling is an indispensable, yet essential step in rendering. In PBRT, we have seen and reviewed several strategies for sampling, including stratified sampling, best-candidate sampling and adaptive sampling.

Poisson-Disk sampling is highly praised for its well-distributed samples, where all samples are at least  $\gamma$  apart without leaving any holes in any regions,  $\gamma$  being a user-determined parameter.

This sort of distribution, often referred to as Blue Noise, is generally considered ideal for applications in rendering. A naïve algorithm to generate Poisson-Disk sampling is the dart-throwing algorithm. However, in dart-throwing, generating random samples and rejecting until a valid one is created might lead to an endless loop and result in undesirable computational cost. Moreover, dart-throwing algorithm potentially under-samples and creates holes in some regions. Many existing algorithms that guarantee better sampling quality in a less time complexity, on the other hand, are not easily generalized to higher dimensions.

Therefore, for applications that require three or more dimensions, e.g., motion blur and depth-of-field, this method proposed by Robert Bridson shed light on generating Blue Noise. The algorithm is guaranteed to run in  $O(N)$  time, where  $N$  is the number of Poisson-Disk samples. We generated candidates the same way as that introduced by Dunbar and Humphreys in 2006, but instead of calculating the allowed scalloped region, we rejected invalid candidates.

## IMPLEMENTATION

We implemented the algorithm in 2D, though it can be easily extended to an arbitrary dimension. The algorithm can be divided into three steps: (i) setting up utility data structures, (ii) initialization for the loop of sample generation, and (iii) the main loop.

### (i) Setting up utility data structures

We first created a  $n$ -dimensional grid, called background grid, in order to accelerate the spatial searches and calculation.  $n$  is the number of dimensions in which we wanted to generate samples, and in our implementation, we set it to two. The size of each cell in the grid is  $\gamma/\sqrt{n}$ , and therefore each cell would contain at most one sample.

### (ii) Initialization for the loop of sample generation


In this step, a randomly created sample is chosen uniformly from our domain, say  $x_0$ . We inserted  $x_0$  into the corresponding cell in the background grid and prepared a container, “*active list*”, as mentioned in the original paper, to accommodate samples. The sample  $x_0$  is then inserted into the active list.

### (iii) The main loop


For the main loop of generating samples, while the active list is not empty, we pop a random sample inside the active list, say  $x_i$ . And then, we generated  $k$  points around  $x_i$  by uniformly choosing from the annulus region between radius  $\gamma$  and  $2\times\gamma$ . For each of the  $k$  points, we used the background grid to check whether it was within distance  $\gamma$  of any existing samples. With background grid, we had to test the nearby samples only, which took constant time rather than linear. If a point is sufficiently far from all existing samples after the test, we inserted it into the background grid as well as active list.

## RESULTS


We demonstrated the sampling algorithm in 4 images. The execution time and number of samples generated are both recorded.




shape: circle  
total # of samples: 9832  
time for generation: 0.316 sec  
(time for drawing: 0.133 sec)



shape: rectangle  
total # of samples: 12480  
time for generation: 0.406 sec  
(time for drawing: 0.227 sec)



shape: circle  
total # of samples: 19607  
time for generation: 0.633 sec  
(time for drawing: 0.183 sec)



shape: rectangle  
total # of samples: 24900  
time for generation: 0.806 sec  
(time for drawing: 0.230 sec)

## USAGE

With the properly composed `Makefile`, compiling and executing the program is simple and easy. Basically there are two main source files, `Poisson.cpp` and `drawPoisson.py`, for generating samples and drawing the results respectively. The following 3 steps help running the program:

1. Use `$make` to compile the C++ program first.

Rendering, 2016 fall

2. Execute with `$make run` command.
3. Lastly, draw the result using `$make draw`.

By default, the C++ program generated samples in a shape of circle, but it can be easily adjusted by setting boolean constant, `Circle`, to `false` in `Poisson.cpp` (line #14). And the constant `k` is set to 30 (line #15) as suggested in the paper.

In addition, to draw the result with white background and black samples, use the command `$make draw-white` instead of `$make draw` in the third step described earlier in this section.

## ENVIRONMENT

Operating System: macOS Sierra (10.12.2)

CPU: 2.6 GHz Intel Core i5

RAM: 8GB 1600 MHz DDR3

GPU: Intel Iris 1536 MB

Disk: 256 GB PCIe SSD

## REFERENCES

[1] Daniel Dunbar, Greg Humphreys. A Spatial Data Structure for Fast Poisson-Disk Sample Generation. Proceedings of SIGGRAPH 2006.

[2] Robert Bridson. Fast Poisson Disk Sampling in Arbitrary Dimensions. Proceedings of SIGGRAPH 2007.